

# Artificial Intelligence

Kushal Shah @ Sitare

# Adversarial Search

- Game Playing [Section 5.1]
- Understand and applying Mini-max [Section 5.2.1]
- Reducing search space with Alpha-Beta Pruning [Section 5.3]

Given an initial state of a 8-puzzle problem and final state to be reached-

2	8	3
1	6	4
7		5

**Initial State**


1	2	3
8		4
7	6	5

**Final State**

Find the most cost-effective path to reach the final state from initial state using A\* Algorithm.


# Games in AI

- Any multi-agent environment is generally a 'game'
- Typical Game components:
  - Well-defined Goal
  - Well-defined Rules
  - Feedback System
  - Voluntary Participation
- In classical AI, the most common games are of a rather specialized kind
  - Deterministic and Fully observable environment
  - Two players taking turns
  - Zero-sum 'adversarial' games of perfect information (eg. chess)

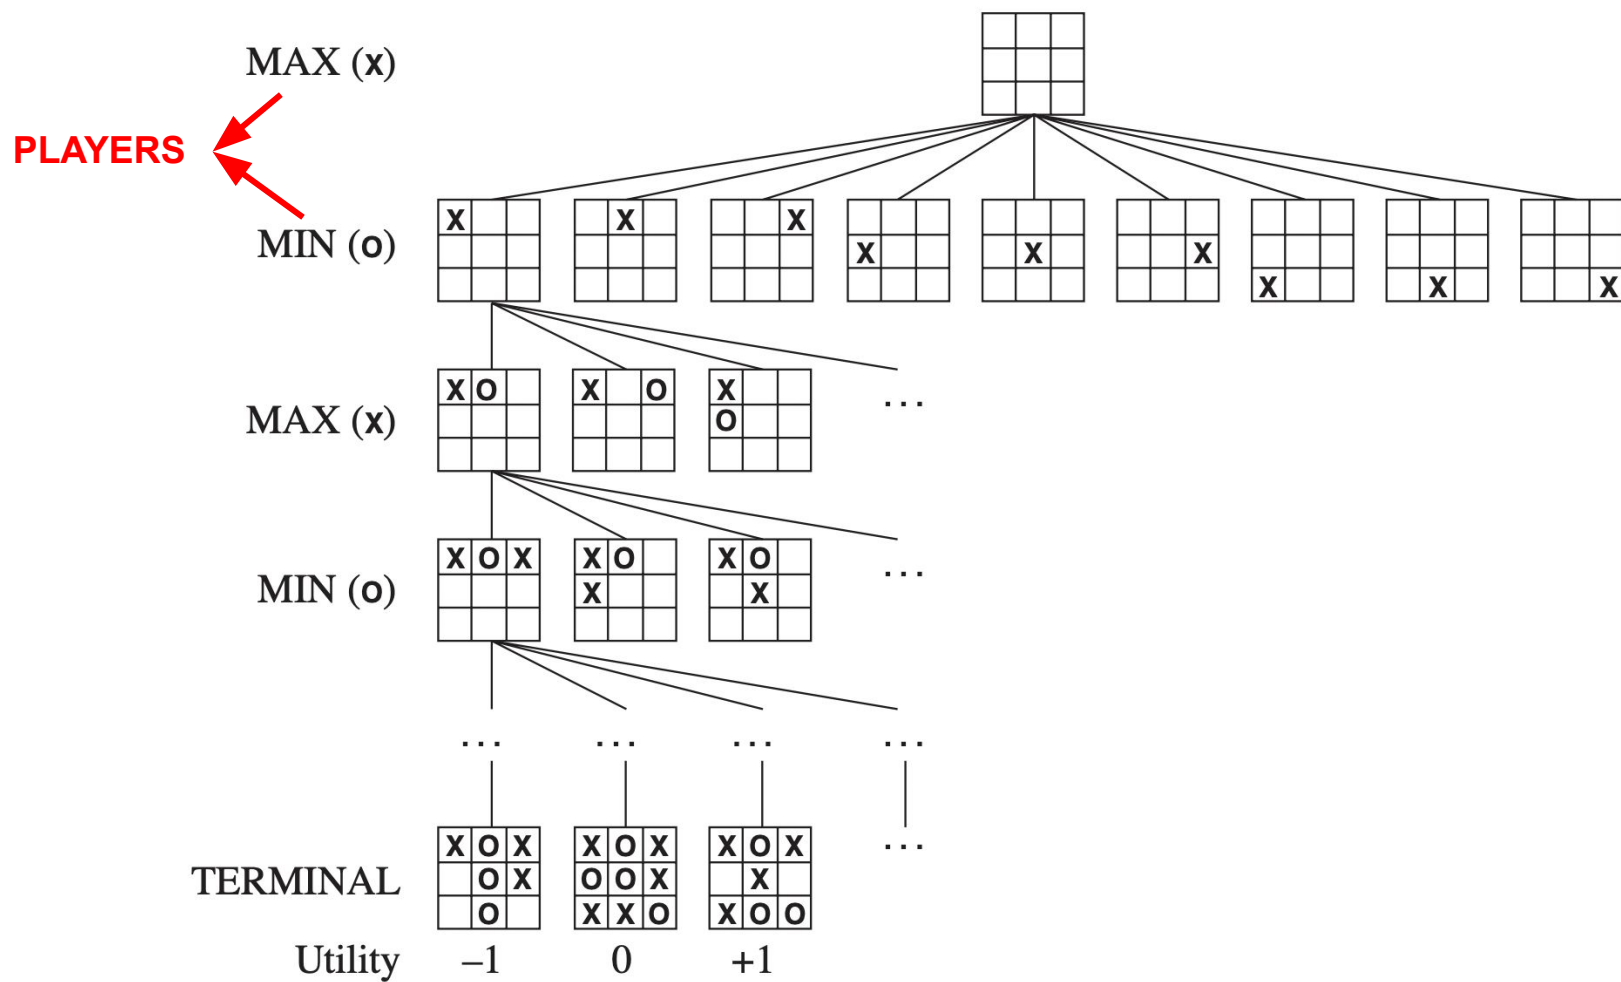


# Games in AI

- Games are interesting in AI because they are too hard to solve algorithmically
  - Humans are so far much better at games in general
- Chess has an average branching factor of about 35, and there are about 50 moves by each player
  - So the search tree has about  $35^{100}$  or  $10^{154}$  nodes
  - The search graph has about  $10^{40}$  distinct nodes
- Games require the ability to make lot of good but non-optimal decisions
- Games also penalize inefficiency severely
  - Take too much time and you lose the game!



**Figure 5.1** A (partial) game tree for the game of tic-tac-toe. The top node is the initial




**Figure 5.1** A (partial) game tree for the game of tic-tac-toe. The top node is the initial





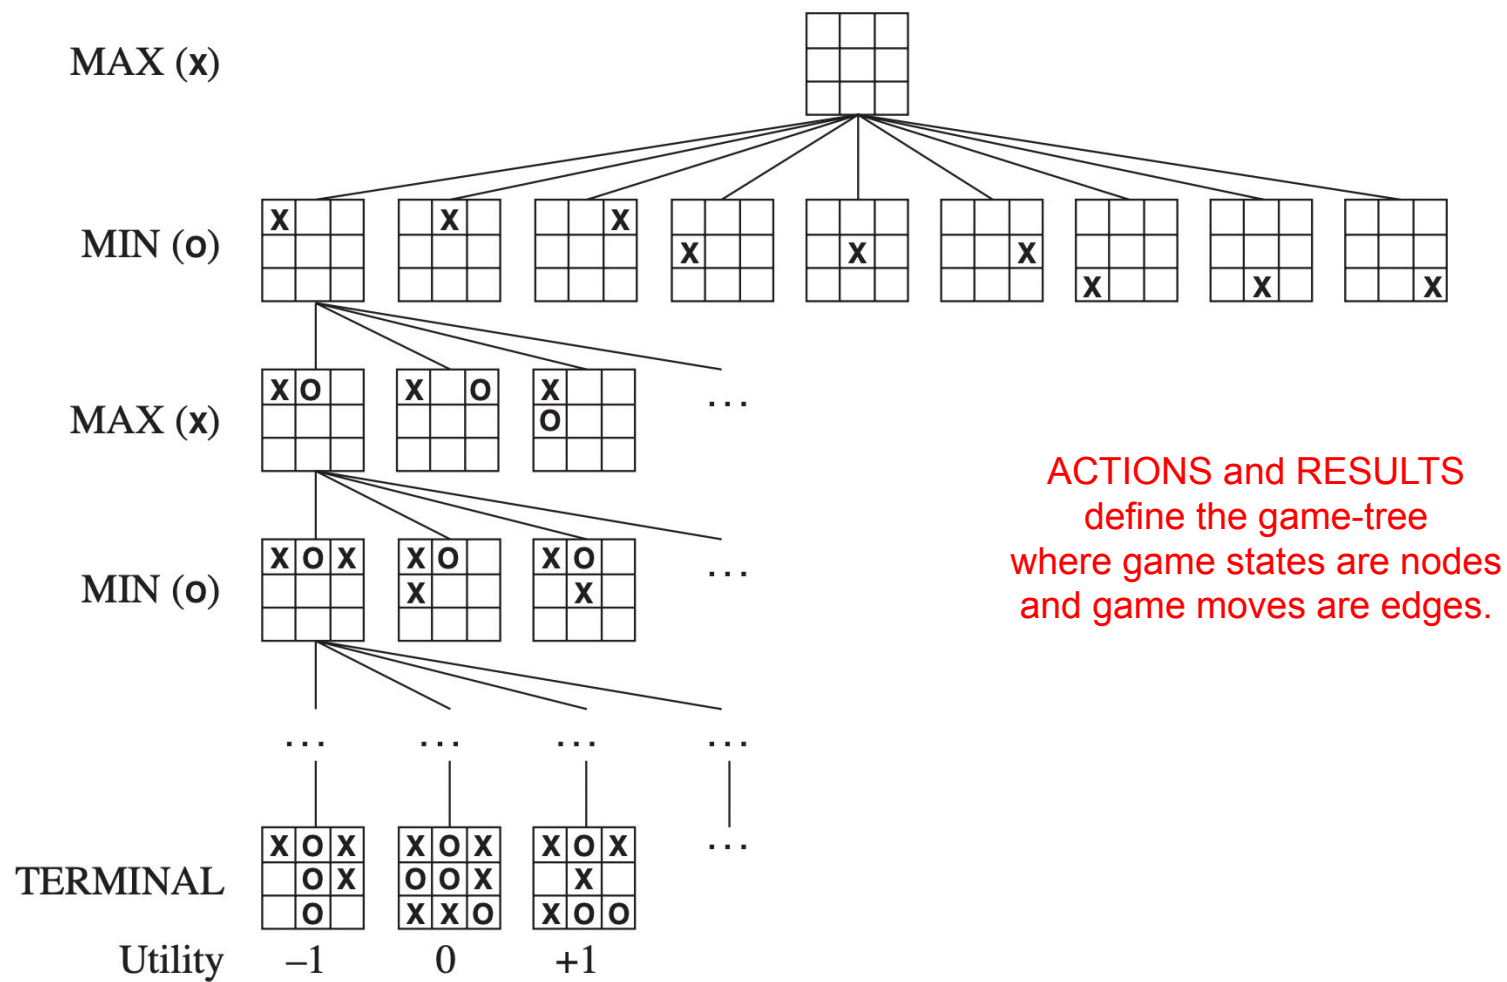




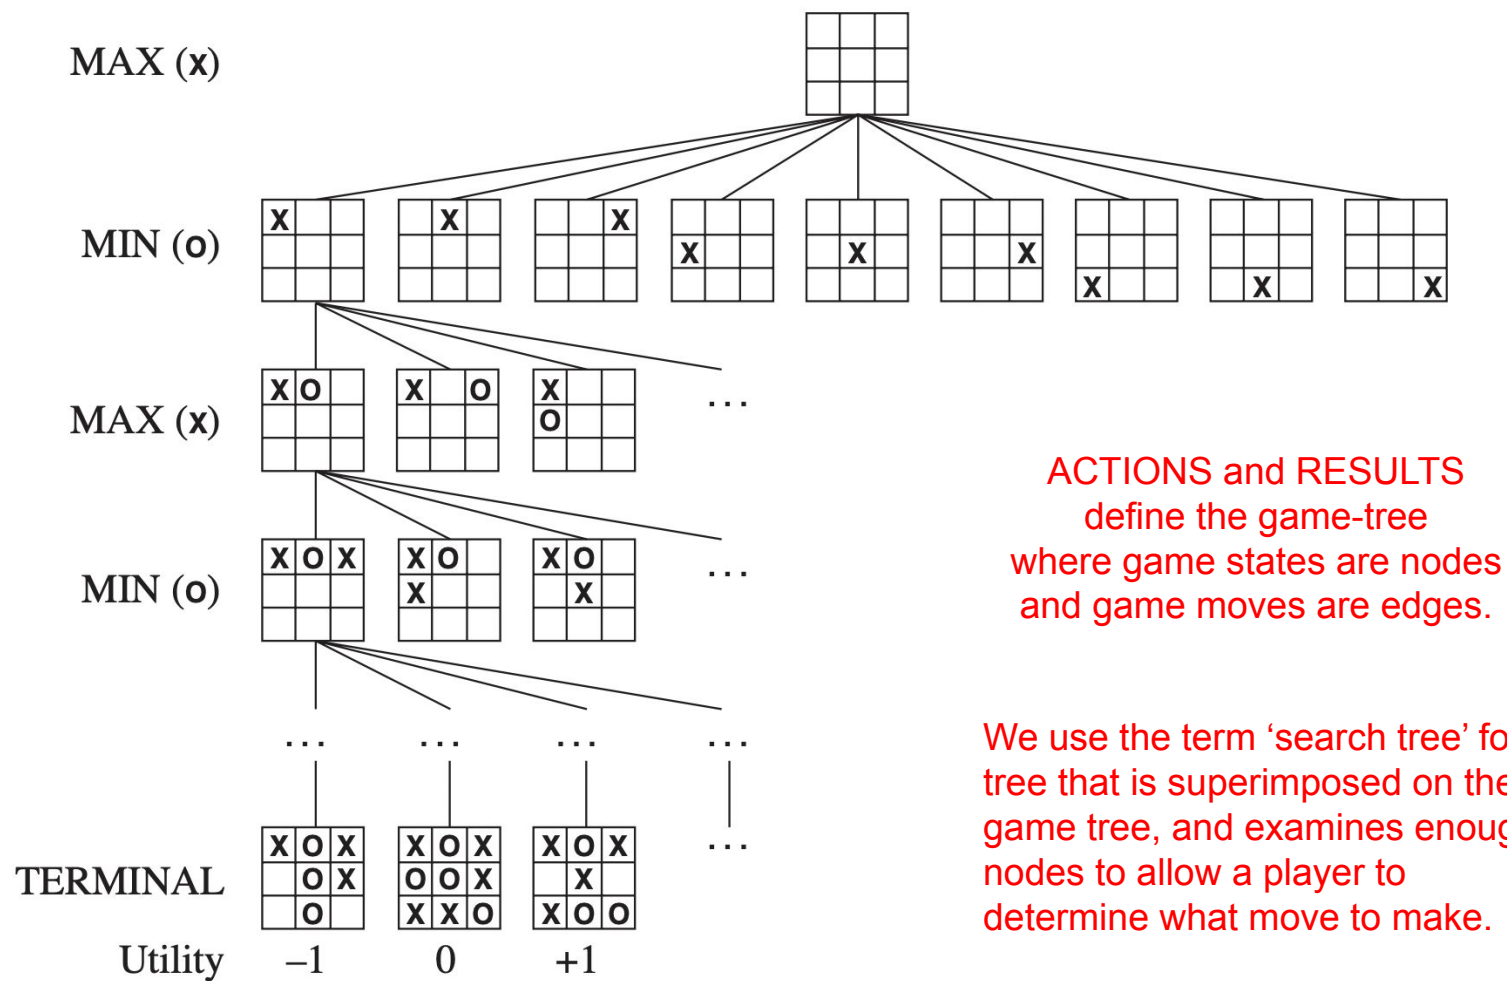
**Figure 5.1** A (partial) game tree for the game of tic-tac-toe. The top node is the initial

# Definition of a AI game

- $S_0$ : The initial state, which specifies how the game is set up at the start.
- PLAYER(s): Defines which player has the move in a state.
- ACTIONS(s): Returns the set of legal moves in a state.
- RESULT(s, a): The transition model, which defines the result of a move.
- TERMINAL-TEST(s): A terminal test, which is true when the game is over and false otherwise
  - States where the game has ended are called terminal states.
- UTILITY(s, p): A utility function (also called an objective function or payoff function), defines the final numeric value for a game that ends in terminal state s for a player p.
  - In chess, the outcome is a win, loss, or draw, with values +1, 0, or  $\frac{1}{2}$
  - The payoffs in backgammon range from 0 to +192
- In a zero-sum game, the total payoff to all players is the same for every instance of the game




**Figure 5.1** A (partial) game tree for the game of tic-tac-toe. The top node is the initial



**Figure 5.1** A (partial) game tree for the game of tic-tac-toe. The top node is the initial

MAX


MIN



**Figure 5.2** A two-ply game tree. The  $\triangle$  nodes are “MAX nodes,” in which it is MAX’s turn to move, and the  $\nabla$  nodes are “MIN nodes.” The terminal nodes show the utility values for MAX; the other nodes are labeled with their minimax values.

MAX


MIN



MiniMax Function

MAX

MIN




## MiniMax Function

- The minimax value of a node is the utility (for MAX) of being in the corresponding state



MAX

MIN




## MiniMax Function

- The minimax value of a node is the utility (for MAX) of being in the corresponding state
- Assume that both players play optimally from there to the end of the game

MAX

MIN




## MiniMax Function

- The minimax value of a node is the utility (for MAX) of being in the corresponding state
- Assume that both players play optimally from there to the end of the game
- The minimax value of a terminal state is just its utility

MAX

MIN




## MiniMax Function

- The minimax value of a node is the utility (for MAX) of being in the corresponding state
- Assume that both players play optimally from there to the end of the game
- The minimax value of a terminal state is just its utility
- Given a choice, MAX prefers to move to a state of maximum value

MAX

MIN




## MiniMax Function

- The minimax value of a node is the utility (for MAX) of being in the corresponding state
- Assume that both players play optimally from there to the end of the game
- The minimax value of a terminal state is just its utility
- Given a choice, MAX prefers to move to a state of maximum value

MAX

MIN




## MiniMax Function

- The minimax value of a node is the utility (for MAX) of being in the corresponding state
- Assume that both players play optimally from there to the end of the game
- The minimax value of a terminal state is just its utility
- Given a choice, MAX prefers to move to a state of maximum value

MAX

MIN




## MiniMax Function

- The minimax value of a node is the utility (for MAX) of being in the corresponding state
- Assume that both players play optimally from there to the end of the game
- The minimax value of a terminal state is just its utility
- Given a choice, MAX prefers to move to a state of maximum value

MAX

MIN




## MiniMax Function

- The minimax value of a node is the utility (for MAX) of being in the corresponding state
- Assume that both players play optimally from there to the end of the game
- The minimax value of a terminal state is just its utility
- Given a choice, MAX prefers to move to a state of maximum value

MAX

MIN




$\text{MINIMAX}(s) =$

{  
|



MAX

MIN




$$\text{MINIMAX}(s) = \begin{cases} \text{UTILITY}(s) \\ \end{cases}$$

if **TERMINAL-TEST**( $s$ )

MAX

MIN




MINIMAX( $s$ ) =

$$\begin{cases} \text{UTILITY}(s) & \text{if } \text{TERMINAL-TEST}(s) \\ \max_{a \in \text{Actions}(s)} \text{MINIMAX}(\text{RESULT}(s, a)) & \text{if } \text{PLAYER}(s) = \text{MAX} \end{cases}$$

MAX

MIN



MINIMAX( $s$ ) =


$$\begin{cases} \text{UTILITY}(s) & \text{if } \text{TERMINAL-TEST}(s) \\ \max_{a \in \text{Actions}(s)} \text{MINIMAX}(\text{RESULT}(s, a)) & \text{if } \text{PLAYER}(s) = \text{MAX} \\ \min_{a \in \text{Actions}(s)} \text{MINIMAX}(\text{RESULT}(s, a)) & \text{if } \text{PLAYER}(s) = \text{MIN} \end{cases}$$

to move  
A

B

C

A




**Figure 5.4** The first three plies of a game tree with three players ( $A$ ,  $B$ ,  $C$ ). Each node is labeled with values from the viewpoint of each player. The best move is marked at the root.

to move  
A

B

C

A




**Figure 5.4** The first three plies of a game tree with three players ( $A$ ,  $B$ ,  $C$ ). Each node is labeled with values from the viewpoint of each player. The best move is marked at the root.

to move  
A

B

C

A




**Figure 5.4** The first three plies of a game tree with three players ( $A$ ,  $B$ ,  $C$ ). Each node is labeled with values from the viewpoint of each player. The best move is marked at the root.

to move  
A

B

C

A




**Figure 5.4** The first three plies of a game tree with three players ( $A$ ,  $B$ ,  $C$ ). Each node is labeled with values from the viewpoint of each player. The best move is marked at the root.

to move  
A

B

C

A




**Figure 5.4** The first three plies of a game tree with three players ( $A$ ,  $B$ ,  $C$ ). Each node is labeled with values from the viewpoint of each player. The best move is marked at the root.



A

C

A




**Figure 5.4** The first three plies of a game tree with three players ( $A$ ,  $B$ ,  $C$ ). Each node is labeled with values from the viewpoint of each player. The best move is marked at the root.

to move  
A

B

C

A




**Figure 5.4** The first three plies of a game tree with three players ( $A$ ,  $B$ ,  $C$ ). Each node is labeled with values from the viewpoint of each player. The best move is marked at the root.

to move  
A

B

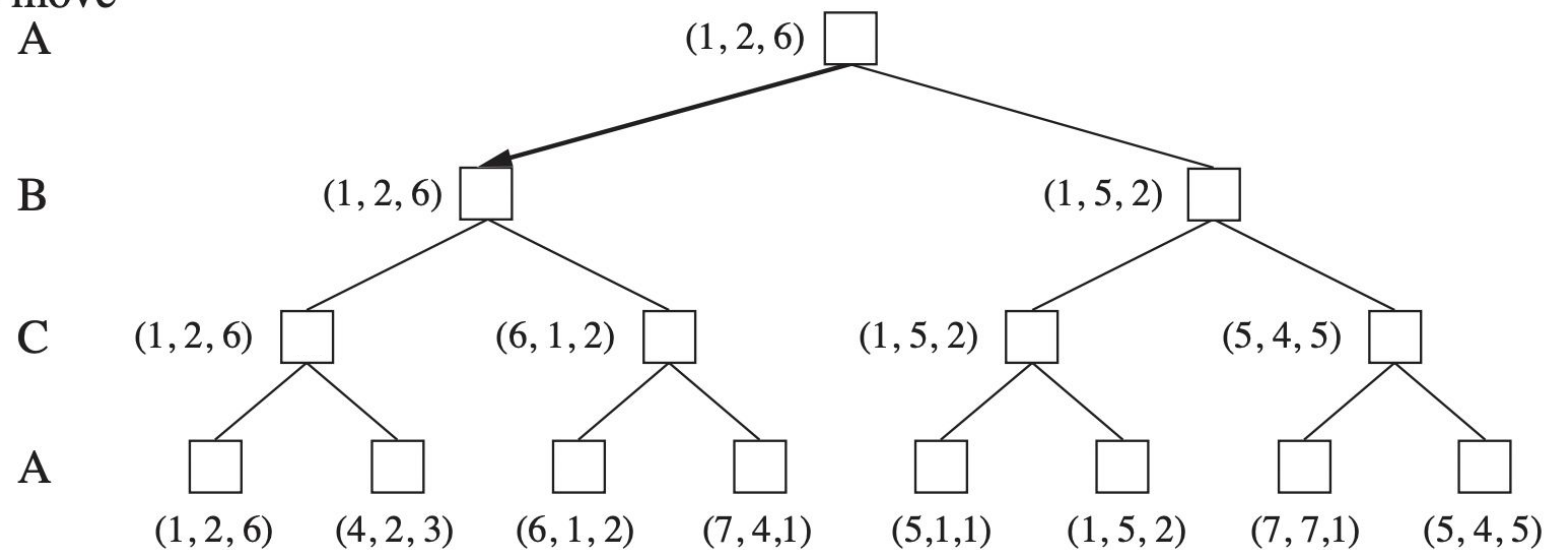
C

A



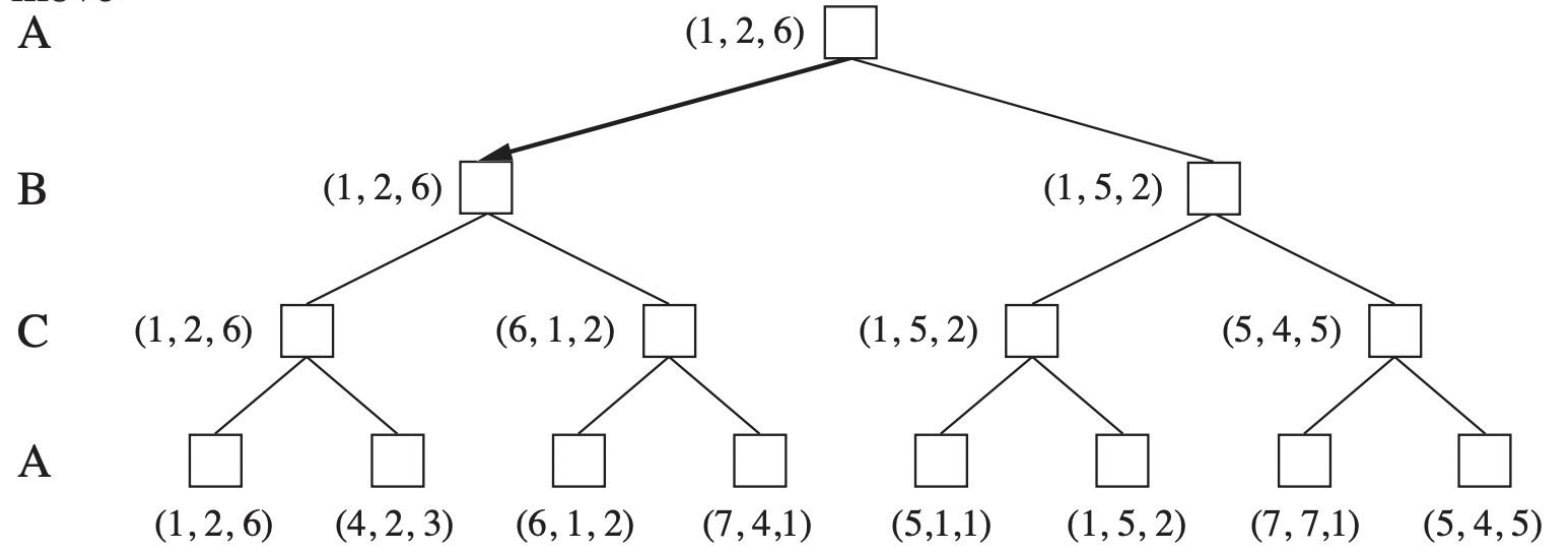
**Figure 5.4** The first three plies of a game tree with three players ( $A$ ,  $B$ ,  $C$ ). Each node is labeled with values from the viewpoint of each player. The best move is marked at the root.

to move  
A



**Figure 5.4** The first three plies of a game tree with three players ( $A$ ,  $B$ ,  $C$ ). Each node is labeled with values from the viewpoint of each player. The best move is marked at the root.

to move  
A



**Figure 5.4** The first three plies of a game tree with three players ( $A$ ,  $B$ ,  $C$ ). Each node is labeled with values from the viewpoint of each player. The best move is marked at the root.

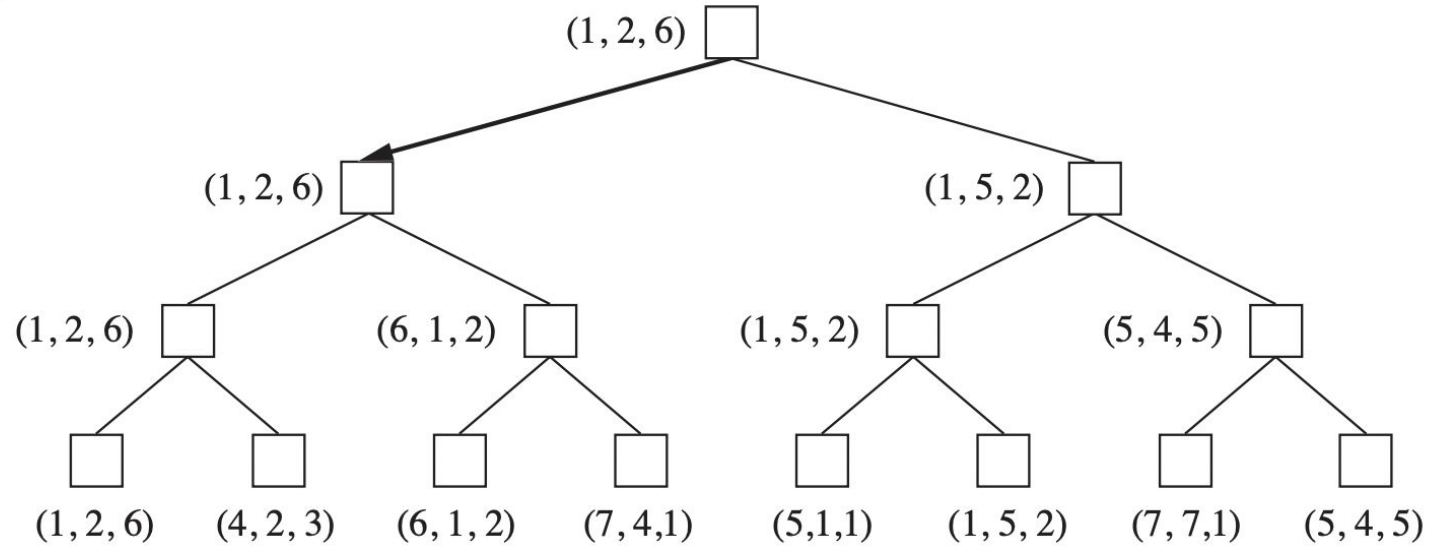
- In a zero-sum three player game, do all the players have to always compete?

to move  
A

B

C

A



**Figure 5.4** The first three plies of a game tree with three players ( $A$ ,  $B$ ,  $C$ ). Each node is labeled with values from the viewpoint of each player. The best move is marked at the root.

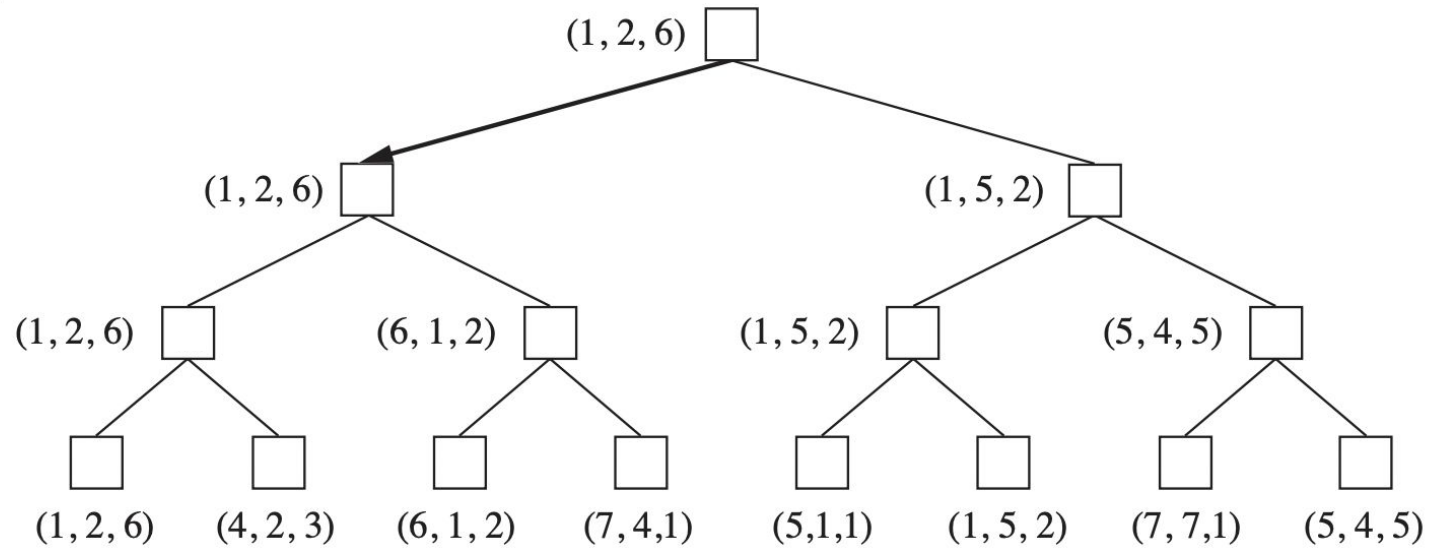
- In a zero-sum three player game, do all the players have to always compete?
- If  $C$  is in a much stronger position, it makes sense for  $A$  and  $B$  to cooperate till they can pull down  $C$

to move  
A

B

C


A




**Figure 5.4** The first three plies of a game tree with three players ( $A$ ,  $B$ ,  $C$ ). Each node is labeled with values from the viewpoint of each player. The best move is marked at the root.

- In a zero-sum three player game, do all the players have to always compete?
- If  $C$  is in a much stronger position, it makes sense for  $A$  and  $B$  to cooperate till they can pull down  $C$  [usually called politics, but its just game play in an organisation or even in families]

What is the MiniMax value at the root node (A) for this 2-player 3-ply game?








Write a python code  
to generate the game tree for Tic-Tac-Toe,  
and apply the MiniMax algorithm.

# ALPHA-BETA PRUNING

MAX


MIN



## MiniMax Function


- The minimax value of a node is the utility (for MAX) of being in the corresponding state
- Assume that both players play optimally from there to the end of the game
- The minimax value of a terminal state is just its utility
- Given a choice, MAX prefers to move to a state of maximum value

(a)




**Figure 5.5** Stages in the calculation of the optimal decision for the game tree in Figure 5.2.



(a)




**Figure 5.5** Stages in the calculation of the optimal decision for the game tree in Figure 5.2.



**Figure 5.5** Stages in the calculation of the optimal decision for the game tree in Figure 5.2.




**Figure 5.5** Stages in the calculation of the optimal decision for the game tree in Figure 5.2.




**Figure 5.5** Stages in the calculation of the optimal decision for the game tree in Figure 5.2.






**Figure 5.5** Stages in the calculation of the optimal decision for the game tree in Figure 5.2.




**Figure 5.5** Stages in the calculation of the optimal decision for the game tree in Figure 5.2.



**Figure 5.5** Stages in the calculation of the optimal decision for the game tree in Figure 5.2.

If we use alpha-beta pruning, which nodes will not be visited by the algorithm?



- IBM Deep Blue was the first machine to beat a reigning world chess champion in a six-game match



One of the two cabinets of Deep Blue in its exhibit at the [Computer History Museum](#), California

<b>Active</b>	1995 (prototype)
	1996 (release)
	1997 (upgrade)



One of the two cabinets of Deep Blue in its exhibit at the [Computer History Museum](#), California

- IBM Deep Blue was the first machine to beat a reigning world chess champion in a six-game match
- Factors that contributed to this success:
  - single-chip chess search engine
  - massively parallel system
  - strong emphasis on search extensions
  - complex evaluation function
    - Alpha-Beta Pruning
    - Iterative Deepening
  - effective use of a Grandmaster game database

**Active**

1995 (prototype)

1996 (release)

1997 (upgrade)



One of the two cabinets of Deep Blue in its exhibit at the [Computer History Museum](#), California

- IBM Deep Blue was the first machine to beat a reigning world chess champion in a six-game match
- Factors that contributed to this success:
  - single-chip chess search engine
  - massively parallel system
  - strong emphasis on search extensions
  - **complex evaluation function**
    - **Alpha-Beta Pruning**
    - **Iterative Deepening**
  - effective use of a Grandmaster game database

<b>Active</b>	1995 (prototype)
	1996 (release)
	1997 (upgrade)



One of the two cabinets of Deep Blue in its exhibit at the [Computer History Museum](#), California

- IBM Deep Blue was the first machine to beat a reigning world chess champion in a six-game match
- Factors that contributed to this success:
  - single-chip chess search engine
  - massively parallel system
  - strong emphasis on search extensions
  - complex evaluation function
    - Alpha-Beta Pruning
    - Iterative Deepening
  - effective use of a Grandmaster game database
- Kasparov demanded a rematch, but IBM had dismantled Deep Blue after its victory and refused the rematch

**Active**

1995 (prototype)

1996 (release)

1997 (upgrade)





One of the two cabinets of Deep Blue in its exhibit at the [Computer History Museum](#), California

- IBM Deep Blue was the first machine to beat a reigning world chess champion in a six-game match
- Factors that contributed to this success:
  - single-chip chess search engine
  - massively parallel system
  - strong emphasis on search extensions
  - complex evaluation function
    - Alpha-Beta Pruning
    - Iterative Deepening
  - effective use of a Grandmaster game database
- Kasparov demanded a rematch, but IBM had dismantled Deep Blue after its victory and refused the rematch
- Inspired the development of IBM Watson for playing Jeopardy

**Active**

1995 (prototype)

1996 (release)

1997 (upgrade)



One of the two cabinets of Deep Blue in its exhibit at the [Computer History Museum](#), California

- IBM Deep Blue was the first machine to beat a reigning world chess champion in a six-game match
- Factors that contributed to this success:
  - single-chip chess search engine
  - massively parallel system
  - strong emphasis on search extensions
  - complex evaluation function
    - Alpha-Beta Pruning
    - Iterative Deepening
  - effective use of a Grandmaster game database
- Kasparov demanded a rematch, but IBM had dismantled Deep Blue after its victory and refused the rematch
- Inspired the development of IBM Watson for playing Jeopardy
- Alpha-beta pruning worked great for Chess, Checkers and Othello

**Active**

1995 (prototype)

1996 (release)

1997 (upgrade)



One of the two cabinets of Deep Blue in its exhibit at the [Computer History Museum](#), California


**Active**

1995 (prototype)

1996 (release)

1997 (upgrade)

- IBM Deep Blue was the first machine to beat a reigning world chess champion in a six-game match
- Factors that contributed to this success:
  - single-chip chess search engine
  - massively parallel system
  - strong emphasis on search extensions
  - complex evaluation function
    - Alpha-Beta Pruning
    - Iterative Deepening
  - effective use of a Grandmaster game database
- Kasparov demanded a rematch, but IBM had dismantled Deep Blue after its victory and refused the rematch
- Inspired the development of IBM Watson for playing Jeopardy
- Alpha-beta pruning worked great for Chess, Checkers and Othello
- AlphaGo uses Monte Carlo Tree Search (MCTS)



Write a python code  
to generate the game tree for Tic-Tac-Toe,  
and apply the Alpha-Beta pruning algorithm.